

Sprint 3 Plan

Team VeriFi · Sprint Completion: July 14, 2026 · Revision 1: 07/08/26

Sprint Period: July 7, 2026 – July 14, 2026

Sprint 3 — Component Integration + Demo

Goal: Integrate all individual components so that the project can be demoed. Users should be able to send a message, and receive documents that would be relevant to them.

User Stories & Tasks

Priorit y	ID	User Story	Story Points
1	3.1	As a fidelity user, I want to be able to see which documents I can refer to in order to answer my question. - Task 1: Finalize contracts between backend and API, and frontend and API - Task 2: Integrate front end requests to receive backend query results - Task 3: Create a clear handoff file for the front end to call - Task 4: Results should display relevant document chunks - Task 5: Handle error states, rate limits, and loading feedback in the chat interface - Task 6: Write integration tests to verify context accuracy Total: 8 story points	8
2	3.2	As an investor using Verifi to research new companies or funds to invest I want to be able to ask financial questions and receive a coherent answer within a reasonable time period. - Task 1: Bridge with backend/AI API to receive and respond to questions. - Task 2: Clean up frontend for easier use - Task 3: Be able to receive the documents related to the question and response Total: 4 story points	3
3	3.3	As a new investor I want to be able to upload documents important to me specifically so that my responses can be more personalized. - Task 1: Build data ingestion pipeline to better handle large documents - Task 2: Experiment to find the best chunk size for context retention Total: 3 story points	3
4	3.4	As a Fidelity user, I want VeriFi to generate answers using only the retrieved document chunks so that I receive accurate responses with supporting citations. - Task 1: Connect the retrieved document chunks from the backend to the prompt builder.	5

Priority	ID	User Story	Story Points
		- Task 2 : Update the prompt builder so it works with the actual chunk field names used by the ingestion and frontend components. - Task 3 : Connect the completed prompt to the LLM response-generation step. - Task 4 : Return the generated answer and its supporting document citations through the /chat endpoint. - Task 5 : Test questions with enough context, insufficient context, and unrelated context. - Task 6 : Validate that each answer is grounded in the retrieved chunks and that its citations match the sources used. Total: 5 story points	

Team Roles

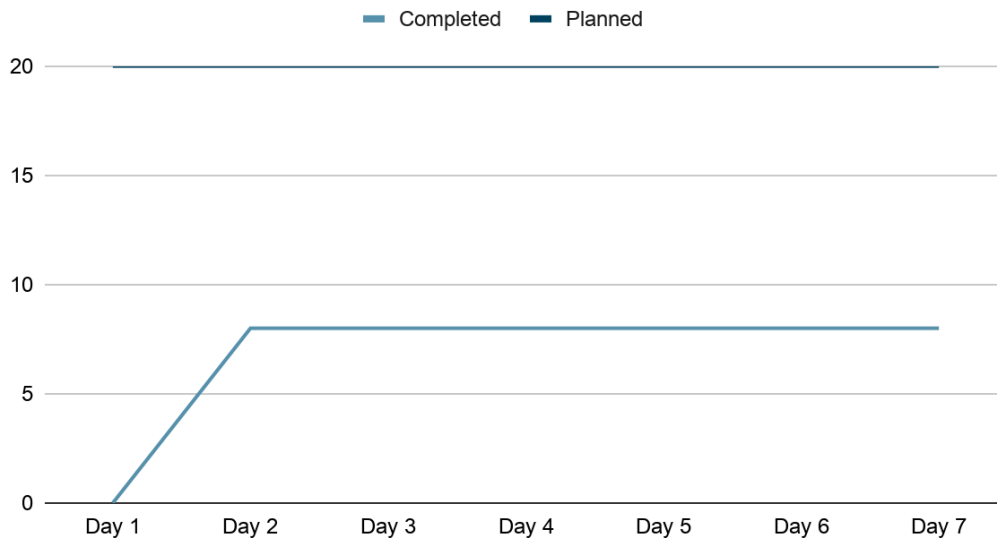
Christopher: Product Owner
 Divya: Developer, Scrum Master
 Ethan: Developer
 Sai: Developer
 Srushti: Developer

Initial Task Assignment

Christopher: 3.4 — Task 6: Validate that each answer is grounded in the retrieved chunks and that its citations match the sources used
 Ethan: 3.1 — First task: Finalize contracts between backend and API, and frontend and API
 Sai: 3.2 — First task: Be able to send and receive basic questions/response from the frontend.
 Divya: 3.4 — First task: Connect retrieved document chunks to the prompt builder and make the data formats consistent across the backend, prompts, and frontend.
 Srushti: 3.3 — First task: Handle multiple files and build data ingestion pipeline to better handle large documents

Initial Burnup Chart

Points scored



Initial Scrum Board

User Stories	Tasks To Do	Tasks In Progress	Tasks Completed
3.1 As a fidelity user, I want to be able to see which documents I can refer to in order to answer my question.			- Finalize API Contract - Accept front end requests and send backend query results - Create clear handoff file for frontend - Results should be relevant chunks and citations - Error handling - Testing
3.2 - As an investor using Verifi to research new companies or funds to invest, I want to be able to ask financial questions and receive a coherent answer within a reasonable time period.	- Bridge with backend/AI API to receive and respond to questions. - Clean up frontend for easier use - Be able to receive the documents related to the question and response		

User Stories	Tasks To Do	Tasks In Progress	Tasks Completed
3.3 - As a new investor I want to be able to upload documents important to me specifically so that my responses can be more personalized.	-		Build data ingestion pipeline to better handle large documents - Experiment to find the best chunk size for context retention
3.4			- Integrated the prompt builder with the backend retrieval pipeline. - Updated the prompt builder and template to use the project's document metadata and improve grounded responses. - Connected prompt generation to the backend and returned responses with supporting citations. - Tested the complete RAG pipeline using multiple policy questions and verified citation accuracy. - Expanded QA testing to validate grounded responses across representative policy questions. - Fixed integration issues and verified compatibility between the frontend, backend, prompt generation, and retrieval components.

Scrum Dates

- Meeting 1: 07/10/2026, 10AM – 10:30AM — With TA
- Meeting 2: 07/09/2026, 3PM - 4PM
- Meeting 3: 07/14/2026, 9AM – 9:30AM — With TA